

# 2.1 Lesson: Multi-Class Classification with Neural Networks

---



## Multi-Class Classification with Neural Networks


Last week, we saw that by using a **sigmoid** activation function in the output layer, we can perform binary or multi-label classification tasks with a neural network. However, a very common task involves classifying a sample into one of a number of different classes.

### Example:

Recognizing handwritten digits from the [MNIST dataset](#)

(<https://www.kaggle.com/datasets/hojjatk/mnist-dataset>) is often considered the “Hello World” of deep learning. The input consists of 28x28 pixel black-and-white images of digits, and the output is a label from 0 to 9.

This form of classification task requires a generalization of the sigmoid function.



**Figure 1: MNIST Dataset**

## The Softmax Function


Recall that the sigmoid turns any floating-point number into a probability between 0 and 1:

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{e^z}{e^z + 1}$$

The **softmax** function extends the sigmoid to cases where there are  $k > 1$  possible classes. It produces a probability distribution over these classes — that is,  $k$  probabilities that sum to 1.0:

$$\text{softmax}(\mathbf{z}) = \left[ \frac{e^{z_1}}{\sum_{i=1}^k e^{z_i}}, \frac{e^{z_2}}{\sum_{i=1}^k e^{z_i}}, \dots, \frac{e^{z_k}}{\sum_{i=1}^k e^{z_i}} \right]$$

The predicted class is simply the one with the highest softmax probability.



**Figure 2: A Multi-Class Classification Network**

## Loss Function for Multi-Class Classification

In this form of classification, each input is assigned to exactly one class out of  $C$  classes. While the loss function is based on **cross-entropy**, how it's computed depends on how we represent the ground truth labels. There are two possibilities.

### Option 1: One-Hot Encoded Labels

Just as with multi-label classification, we could encode the target as a one-hot vector of 0's and 1's. For example, the ground truth for the diagram above might be:

$$y^{(i)} = [y_0^{(i)}, y_1^{(i)}, y_2^{(i)}] = [0, 0, 1].$$

In this case, the **categorical cross-entropy** loss is computed as:

$$L = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^C y_j^{(i)} \log(\hat{y}_j^{(i)})$$

where:

- $y_j^{(i)} \in \{0,1\}$  : 1 if class  $j$  is correct for sample  $i$ , 0 otherwise.
- $\hat{y}_j^{(i)} \in [0..1]$  : softmax predicted probability for class  $j$ .

Notice how this formula generalizes the binary cross-entropy loss formula, where  $y^{(i)}$  is either 0 or 1 and selects the appropriate part of the sum, to the one-hot vector, where exactly one of the  $y_j^{(i)} = 1$ .

### Option 2: Integer Class Labels

Instead of one-hot encoding, we can simply store the true class as an **integer** in the range  $[0, \dots, C-1]$ . The ground truth for the above diagram would simply be 2. This approach is more efficient, widely used in practice, and is the convention in Keras/TensorFlow.

Now the loss, called Sparse Cross-Entropy Loss, simplifies to:

$$L = -\frac{1}{n} \sum_{i=1}^n \log(\hat{y}_{y^{(i)}}^{(i)})$$

where:

- $y^{(i)}$  is the correct class label for sample  $i$ , used here as an index selecting the appropriate prediction.
- $\hat{y}_{y^{(i)}}^{(i)}$  is the predicted probability (from softmax) for that class.

In this case, the subscript  $y^{(i)}$  indicates which class the prediction must match and ignores all the other labels.